

John Doe

✉ john.doe@example.com | ☎ +49 163 555 1584 | 🐙 github.com/johndoe | 💼 linkedin.com/in/johndoe
| 📍 Berlin, Germany

September 16, 2026

To the Nexus AI Platform Team:

Nexus AI's challenge of serving over 10 billion predictions per day at sub-50ms latency is exactly the kind of infrastructure problem I find most compelling. At Lumino AI, I have spent the last three years building and operating the model serving layer from the ground up — and I would like to bring that experience to your Platform team.

The work I am most proud of maps directly to the core responsibilities of this role. I designed and shipped a real-time feature store (Redis + PostgreSQL) that eliminated ~60% of duplicate feature computation across three model teams, and I led the inference infrastructure overhaul that reduced p99 latency by 40% through ONNX quantisation and dynamic batching. At DataStream, I built the ETL backbone for a 12-source data platform using FastAPI and Celery, with idempotent pipelines that cut failure rates by 35% — which taught me that production data infrastructure is as much about failure modes as the happy path.

Nexus AI's scale is a step beyond where I have operated, and that is precisely why I am interested. I am confident the fundamentals transfer, and I am motivated by the engineering challenges that come with another order of magnitude. I would welcome the chance to discuss how my experience maps to what your team is building.

Sincerely,

John Doe